Introduction to Modeling: Takeaways ⤷

Syntax

VISUALIZING BIVARIATE RELATIONSHIPS

- Generating scatterplots to visualize bivariate relationships:

```
ggplot(data = uber_trips,

    aes(x = distance, y = cost)) +

    geom_point()
```

- Visualizing a linear regression line:

```
ggplot(data = uber_trips,

    aes(x = distance, y = cost)) +

    geom_point() +

    geom_smooth(method = "lm", se = FALSE)
```

ANALYZING THE RESIDUALS

- Calculating mean abolute error (MAE):

```
MAE <- mean(abs(df$residuals))
```

Notation

GENERAL FORM OF A PREDICTIVE MODEL:

- In this context,   represents a set of **inputs** and   represents a set of **outputs**. The random error term   is independent of   and has a mean of approximately zero. In reality, the error term is unknown, so we can represent our estimate of   as a function of   as:

- We can omit the error term because it averages to zero. The "hat" symbol indicates an estimate.

BIVARIATE LINEAR MODEL:

- Here   indicates a prediction of   assuming   . Intercept (   ) refers to the value on the y-axis where the value on the x-axis is equal to zero. Slope (   )is the change in   for every single unit change in   .

MEAN ABSOLUTE ERROR (MAE):

- = observed output value

- = predicted output value

-   = the sum of…

-     | = the absolute value of each residual

- = divide the above by the total number of observations (to return the average value)

Concepts

- An *input, or* input variable *is also sometimes referred to as a* predictor, independent variable, feature, attribute, descriptor, *or simply* variable*.

- An *output* or *output variable* is also known as a *dependent variable*, *outcome*, *response variable*, *response*, *target*, or *class.*

- **Prediction:** If our primary purpose of building a model is to generate an accurate prediction, we aren't too concerned if the function of the input to predict the response is a "black box." Rather than understanding , our primary concern is that our model gives us accurate output predictions for each input.

- **Inference:** When motivated by inference, we may or may not be interested in generating predictions for our response. Instead, we wish to understand and how the response variable is affected by changes in the input variable.

- **Error:** The accuracy of our prediction for an output depends on two types of error: *reducible error*, and *irreducible error.* Even though the "E" in mean absolute error stands for error, it does not refer to the epsilon error described here.

- **Reducible error** can be minimized by selecting a statistical method that provides a good estimate of the response. In linear regression, one key to minimizing reducible error is to select the input variable that provides the most accurate estimate of the output vatiable.

- **Irreducible error** is out of our control. Unmeasurable variation contributes to error. Another term for this is random noise.

- **Residual:** The difference between the observed value and the model's prediction.

- **Summary measures:** Statisticians have developed various summary measurements (mean absolute error, for example) that can take the residuals from our model and transform them into a single value that represents the predictive ability of our model.

- **Overfitting:** When a model finds patterns in the training data that are not present in the unseen data.

Resources

- [An Introduction to Statistical Learning with Applications in R by James et al.](#)
- [Applied Predictive Modeling by Max Kuhn and Kjell Johnson.](#)